

Master Project Report

August 30, 2016

GLAPP

Fabio Isler (09-115-965)
Man Wai Li (14-705-156)
Dinesh Pothineni (14-707-988)
Riccardo Patane (XX-XXX-XXX)

supervised by

Prof. Dr. Harald C. Gall

Dr. Philipp Leitner



**University of
Zurich** ^{UZH}



Abstract

Abstract

4 PAGES: Introduction

TODO ▷*DONE BY: Fabio*◁

Cloud computing has created a paradigm shift in the last few years, by making infrastructure available at lower costs and with higher efficiency of operations. These solutions are increasingly being adopted by enterprises and developers, as they can provision huge amount of resources to scale on-demand in order to meet their business needs. In cloud computing, resources such as CPU processing time, disk space, or networking capabilities, are rented and released as a service. Today, the most important model for delivering on the cloud promise is the Infrastructure-as-a-Service (IaaS) paradigm. In IaaS, virtual computing resources are acquired and released on demand, either via an Application Programming Interface (API) or web interface. Along with great flexibility of being able to get new resources on demand and pay for what you use, new problems arise. Selecting a cloud service provider can often be a quite challenging decision for the developer or a company, so is being able to monitor and evaluate these infrastructure resources on a regular basis. With so much of variance in cost and performance, it is imperative that one would look for a reliable deployment, monitoring solutions to strike a balance with application requirements. Furthermore, the complexity and skill required to manage multiple layers of application, data, middleware and operating systems can .

Understanding run time performance, behavior of various application components in real-time can enable us to take advantage of arbitrage opportunities that exist between different machines/regions or even other cloud providers. For example, an application can take advantage by moving closer to its users based on timezones or traffic to improve response time. Currently there is no flexibility to move freely between various cloud providers without great development effort and cost, however such an ability to move freely between providers enables us to benefit from cost and performance differences. We propose a cloud middleware that can take not only take care of application deployment to cloud, but also constantly monitor and trigger such necessary adaptations to benefit from these opportunities. Ideally this middleware also enables developers to specify their intended goals in terms of high level policies to govern the application behavior. The middleware can then break down these policies into low level objectives, in order to trigger adaptations by changing the state of application when required.

1.1 Global Living Cloud Applications

The aim of this project is to develop a middleware for what we call "global living cloud applications" (GLAs). In a nutshell, GLAs are a bio-inspired notion of cloud-native applications. GLAs "live" in the cloud, and are able to migrate between data centers and cloud providers automatically, based on changes in cost and performance of cloud offerings, changes in customer behavior or requirements, or other factors.

TODO ▷*"What I would really do is invest more thought and discussions on the conceptual model behind GLAs - I think there is something pretty cool here, but it does not seem very well-developed to me."*◁

The bio-inspired terminology applies for the different levels of components of a GLA:

- Cell: A cell is the lowest-level component of a GLA, consisting of the actual processes `TODO` ▷ "I wonder whether it makes sense to equate cells to "processes". Is that really the same?"◁
- Organ: An organ consists of one or more cells of the same type `TODO` ▷ "For organs, I understand that organs typically also have cells of different type, i.e., an injected load balancer. I would try to make the model clear in this regard. For instance, it may make sense to distinguish between cells that are user-defined and those that come from your middleware."◁
- GLA: The GLA itself is a collection of organs that form the whole application

In order to manage these GLAs, we introduce a middleware called GLAPP (Global Living Application Platform). It allows a developer to deploy multiple GLAs on whatever cloud he/she has access to and sets a centralized mechanism in place to constantly monitor and manage all the GLAs. The middleware supports heterogeneous environment a GLA can live and move across different providers, regions, instance types, etc. GLAPP is an open source middleware to be used on a private/individual basis. This means that in order to be able to deploy a GLA through GLAPP, it is first required to install the GLAPP platform on the own infrastructure. But no matter where the platform is installed, all GLAs will live in the cloud.

6 PAGES: Background & Architecture

2.1 Basic Design Decisions

TODO ▷ *IN CHARGE: Adrian*◁

2.1.1 Main Components: Provisioning Backend, Frontend, MAPE

TODO ▷ *Describe the design decisions of the main components*◁

2.1.2 Deployment: Containerization with Docker

Containerization vs. other virtualization methods:

TODO ▷ *Explain the advantage of containerization compared to e.g. virtual machines*◁

Docker vs. other containerization implementations:

TODO ▷ *Explain the advantage of Docker compared to e.g. OpenVZ*◁

2.1.3 Orchestration: Docker Swarm

TODO ▷ *Explain the advantage of Docker Swarm compared to e.g. Kubernetes*◁

2.1.4 Rule-Based Adaptations vs. Markov Decision Process

TODO ▷ *DONE BY: Adrian / Riccardo*◁ **TODO** ▷ *Explain the thought process behind this decision*◁

2.2 Implementation Decisions

2.2.1 Backend: Node.js

The backend has been implemented with the *Node.js*¹ runtime. More specifically, we used a framework called *Sails*². Although it is labeled as a complete MVC framework, we mainly used

¹<https://nodejs.org/>

²<http://sailsjs.org/>

it to implement the server-side functionality and to provide a clean REST for the other GLAPP components.

There are two reasons why we chose Node.js and Sails in particular. For one thing, it was the backend framework some of us were most experienced in, which made it easier to focus on the actual project without having to work our way into a new framework. For another thing, Node.js comes with the well-established packaging system *npm*³, where we found very helpful packages for our problem domain. As an example, *dockerode*⁴ was the module which we used to implement the Docker-specific functionality, and it proved to be very well maintained, adding new functionalities whenever a new Docker release was published. Since we were about to work with relatively new technologies like Docker, it seemed reasonable to count on packages which we knew will be adapted to changes in these new technologies.

2.2.2 GUI: AngularJS

TODO ▷ *DONE BY: Dinesh* **TODO** ▷ *Explain the choice and alternatives*

2.2.3 MAPE: Java

TODO ▷ *DONE BY: Adrian / Riccardo* **TODO** ▷ *Explain the choice and alternatives*

2.3 Architecture

The platform consists of 3 different parts/blocks: a front end, a provisioning component and a control loop. The front end provides an interface for developer to interact with the middleware to deploy and manage his/her GLAs. The provisioning component provides the management functionalities of the middleware including cloud infrastructure management, application deployment and access to the application status information. Lastly the control loop is the component responsible for enabling the management of GLAs by the middleware itself. It follows the MAPE (Monitoring, Analysis, Planning and Execution) principle. Possible execution actions are moving cells between different cloud instances (migration), duplicating/splitting cells of the GLA (mitosis), or removing cells.

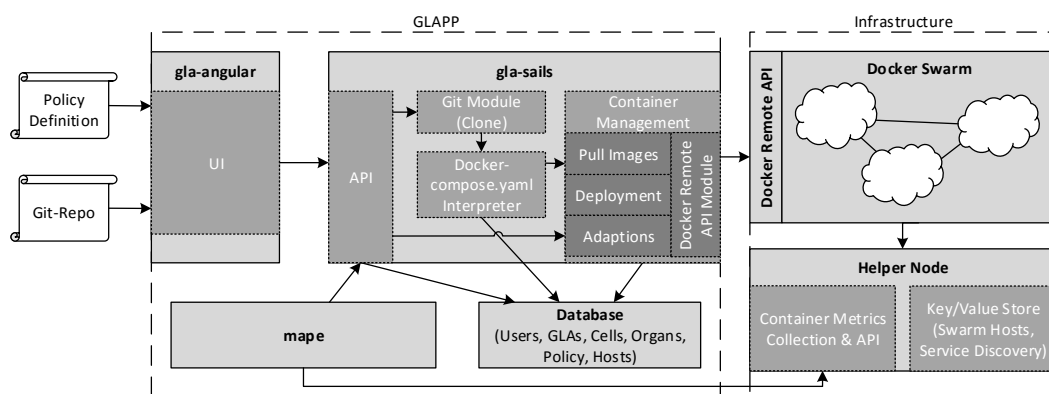


Figure 2.1: Detailed Architecture of GLAPP and the Infrastructure

TODO ▷ *Explain the separate parts - DONE BY: Dinesh*

³<https://www.npmjs.com/>

⁴<https://www.npmjs.com/package/dockerode>

TODO ▷ "I like Fig2 better, but there is still plenty of room for improvement. For instance, I would remove the word "optimal" from the figure - especially, if we now use this iterative Q-learning approach, our deployment will often **not** be optimal, at least temporarily. I am also quite suspicious that the msot dominant block in the figure is the provisioner. Is this really what we want to emphasize?"◁

10 PAGES: Components

3.1 Backend

The backend is mainly responsible for the deployment and is the central component that defines the data structure of the platform. In order to understand the main functionality of the backend server, this section introduces the data models, the REST API, the services and some additional information about this component.

3.1.1 Data Models

Figure 3.1 provides a simple ER-like overview over the way how the server stores the data about users, applications, etc. The mentioned attributes are exemplary, but represent the most important ones to understand the whole model. It is apparent that the model is tightly coupled to the introduced notion of a GLA - an application which consists of organs, which again consist of cells.

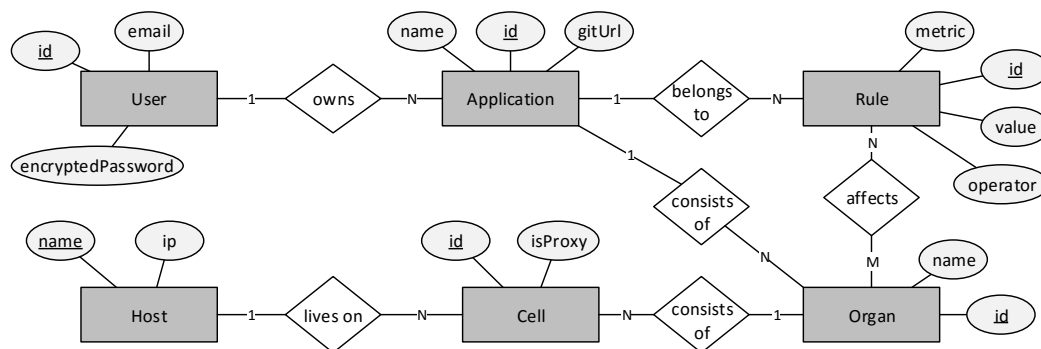


Figure 3.1: Simple overview over the models and their relations

Sails is built upon an ORM called *waterline*¹, which allowed us to manage the data on a relatively high level. That means that by defining the models and their relations and by setting up a connection to any database in the configuration of Sails, the data could be stored and retrieved by simple function calls. We decided to use *MongoDB*² as a database.

¹<https://github.com/balderdashy/waterline>

²<https://www.mongodb.com/community>

3.1.2 REST API

The backend is supposed to serve the frontend and the MAPE component with needed data and to provide an interface for triggering adaptations to the deployment. To achieve this clean interface, there is a REST API in place. Table 3.1 lists and describes the endpoints of this API. Since our components are completely modularized and the communication only happens via this interface, the remaining components like the frontend or the MAPE are easily interchangeable.

3.1.3 Services

In the Sails framework, services are used to handle specific, API-independent functionalities. For most of the internal logic apart from the actual API of the backend server, we created the following two well-encapsulated services:

- **DockerService:**

The DockerService is responsible for all the low-level, docker-specific functionalities. Many of the requests mentioned in table 3.1 trigger quite complex actions. For example, when a user requests to deploy an application, the server generally handles the following steps: (1) Creating an overlay network on the hosts, (2) Creating all the necessary cells in this network, and (3) Creating a reverse proxy cell for each organ that has an exposed port.

To avoid such low-level logic in the API controllers, it makes sense to encapsulate this in a service. The DockerService handles the connection to the Docker infrastructure by the mentioned npm package *dockerode*. It takes care of everything that is concerned with Docker, consisting of the following categories:

- Initialization: Establishing the Docker infrastructure connection, extracting the organs from the initial docker-compose.yml file
- Deployment: Handle the overlay network, deployment (creating the containers and the reverse proxies as cells)
- Adaptations: Moving containers, adding/removing containers, clearing all containers of a certain application
- Infrastructure: Get host information and parse it in order to save it in our database

- **PrometheusService:**

The PrometheusService is responsible for a clean provisioning of analytics data for the frontend. It establishes the connection to the Prometheus API and allows the fetching of any data the user requests. Since this data is cell-level data, it makes sense to have aggregation functions in place. In the current version of this service, only *average* is implemented, but it's programmed in a pipeline-like way so various alterations and aggregations of the data can be added and chained to the service call.

- Dockerfile & docker-compose
- Tests

3.2 2 PAGES: GUI

TODO ▷ *DONE BY: Dinesh*◁

3.3 4 PAGES: MAPE

TODO ▷ *DONE BY: Adrian / Riccardo*◁ Being the control loop of the whole platform, MAPE interfaces with SAILS to retrieve environment information, with monitoring system to retrieve

Table 3.1: Endpoints of the REST API

Endpoint	Description
User	
<i>POST /user/signup</i>	Used by the frontend to sign up a user.
<i>PUT /user/login</i>	Used by the frontend to log in a user.
<i>GET /user/logout</i>	Used by the frontend to log out a user.
<i>GET /user/confirm-login</i>	Helper endpoint needed for the frontend to assure a user stays logged in on a refresh of the page.
Application	
<i>GET /application/getUserApps</i>	Used by the frontend to get high-level information about applications of a specific user.
<i>GET /application/details</i>	Used by the frontend to get detailed information of a specific application.
<i>POST /application/add</i>	Used by the frontend to add an application.
<i>POST /application/remove</i>	Used by the frontend to remove an application, including clearing up its existing cells on the hosts.
<i>POST /application/deploy</i>	Endpoint used by the frontend to deploy an application on the infrastructure.
<i>POST /application/undeploy</i>	Used by the frontend to clear up an application's existing cells on the hosts.
<i>POST /application/rename</i>	Used by the frontend to rename an application.
Organ	
<i>POST /organ/scaleUp</i>	Used by the MAPE to scale up an organ. Constraints can be defined to place the new cell on a specific host, e.g. in the region "US".
<i>POST /organ/scaleDown</i>	Used by the MAPE to scale down an organ. A specific cell can be specified to control what cell is going to be destroyed.
Cell	
<i>POST /cell/move</i>	Used by the MAPE to move a specific cell from one host to another.
Host	
<i>GET /host/info</i>	Used by the MAPE and the frontend get information about the current infrastructure.
<i>GET /host/prometheusUrl</i>	Used by the MAPE get obtain the URL of Prometheus.
Rule	
<i>GET /policy</i>	Used by the MAPE and the frontend to get the list of rules for a specific application.
<i>POST /policy/set</i>	Used by the frontend to set a new or replace an existing rule.
<i>POST /policy/remove</i>	Used by the frontend to remove a rule.
Analytics	
<i>GET /analytics/organCpu</i>	Used by the frontend to retrieve organ-level data for the cpu usage graph.
<i>GET /analytics/organMemory</i>	Used by the frontend to retrieve organ-level data for the memory usage graph.
<i>GET /analytics/events</i>	Used by the frontend to retrieve application-level data for the events graph.

performance metrics, perform computation within itself to decide adaptation action and, in the end, communicate to SAILS to execute the adaptation action.

Regarding the interface with SAILS and monitoring system, HTTP protocol is used with JSON objects in the request and response. The use of JSON objects allows an data structure that is flexible to represent different types of information and to facilitate data exchange. The information retrieved from SAILS includes the infrastructure information of virtual machines, deployment information of application components and their underlying containers. It also includes user-defined policy from the database as well as performance metrics from the backend monitoring system.

Prometheus TODO ▷add reference◁ is a monitoring system that can provide both infrastructure and application performance metrics. Infrastructure metrics are provided by exporters, which is a set of library that expose metrics data of the host environments and the Docker containers. Application metrics are provided with the use of custom-built program to expose the needed data in a Prometheus exposition format TODO ▷add reference◁. The use of Prometheus and its exposition format enable MAPE to access custom metrics such as application performance data or cost metrics of the host from various cloud service providers.

TODO ▷elaborate about healthiness computation◁

Prometheus is also sophisticated in providing different types of query. It supports range query and point query with additional functions that can be applied on the raw data when formulating a query. This enables the fine-grained extraction of data that best fits the corresponding performance metric. For instance, the rate function allows the direct retrieval of the rate of change in CPU time of a container with configurable data point range, data point interval and duration covered by each data point. The robust query functionality provides two major advantages. First of all, it reduces the complexity of computation logic of MAPE that is responsible for getting specific performance metric for subsequent healthiness computation. Secondly, it allows quick tweaking of query parameters to obtain the most relevant data for adaptation action determination during later development stage.

After the healthiness values of the each individual cell is computed based on the policy and performance metric, all violating cells within an organ are taken into account for an evaluation. When the number of violating cell exceed a defined threshold, the organ will be considered as unhealthy. Once an organ is unhealthy, information of all violations will be further processed in MAPE to determine appropriate adaption action.

3.4 2 PAGES: Other

3.4.1 Monitoring / Prometheus

TODO ▷DONE BY: Adrian / Riccardo◁

3.4.2 Supporting Applications

TODO ▷Infrastructure Scripts, Consul, Registrar, Voting-App, Proxy, metrics-server etc.◁ TODO ▷DONE BY: Fabio◁

7 PAGES: Case Study

TODO ▷ *DONE BY: Riccardo* ◁

4.1 Explain Demo Application

- Explain concept of the app at high level
- Explain the components themselves
- Explain the docker-compose file
- Explain our custom metrics (costs, clicks)
- Explain modifications to the app (region switch button, POST request with click-origin to metrics-server with every click)

4.2 Scenarios

- Explain the scenarios
- Explain the triggers
- Explain how MAPE realizes a necessary adaptation

2 PAGES: Future Work

TODO ▷ *DONE BY: Fabio / Adrian*◁

1 PAGE: Conclusion

Bibliography